

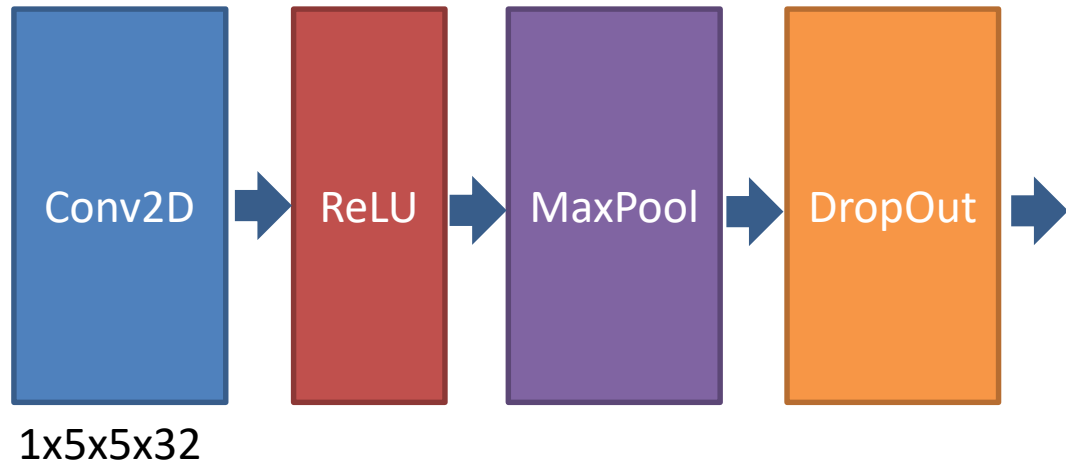
# Kaggle Task: Digit Recognizer Presentation

Xia Rui

# Resources

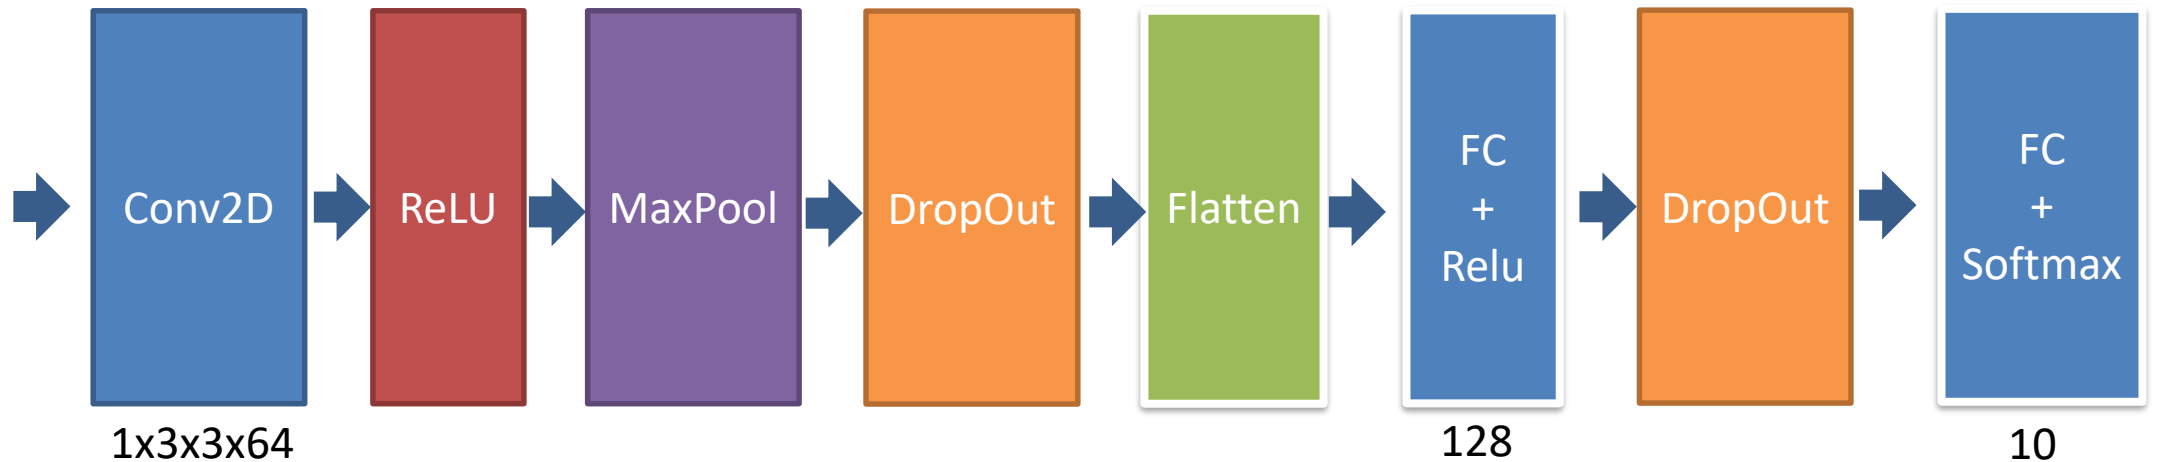
- Stanford University cs231n Convolutional Neural Networks for Visual Recognition
- Python / Jupyter Notebook / Keras
- Research Articles on Batch Normalization / softmax + crossentropy
- Other's kernel

# Basic Models



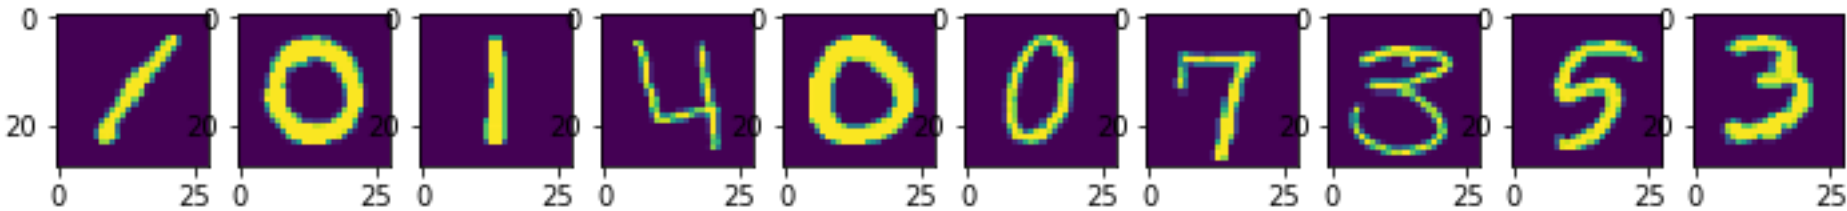
Idea:

1. Leaky ReLU
2. Add more Conv2D layers
3. Batch Normalization
4. Increase epochs
5. Data argumentation



# Data Preparation

label	pixel 0	pixel 1	pixel 2	pixel 3	pixel 4	pixel 5	pixel 6	pixel 7	pixel 8	...	pixel 773	pixel 774	pixel 775	pixel 776	pixel 777	pixel 778	pixel 779	pixel 780	pixel 781	pixel 782	pixel 783
1	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0

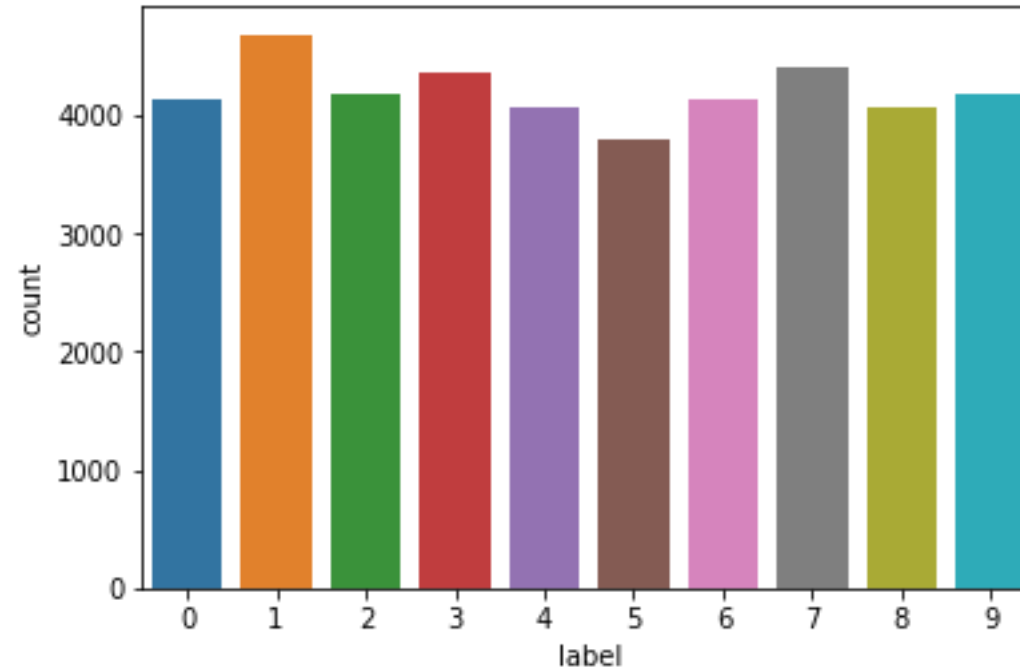


784 Pixels = 28x28 Pixels    integer 0-255    10 digits: from 0-9

# Data Preparation

- Data Distribution
- Checking for Missing Values
- Normalization /255
- To reduce the illumination's influence
- Gradient Descent will converge faster on input range (0,1) than (0,255)

$$\frac{x - x_{min}}{x_{max} - x_{min}}$$



- Data reshape 1x784 to 28x28
- Encode labels to one hot vectors 9 -> (0,0,0,0,0,0,0,0,0,1)
- Convert a class vector (integers) to binary class matrix
- Split the training set to training set and validation set for fitting

# Convolutional Layers



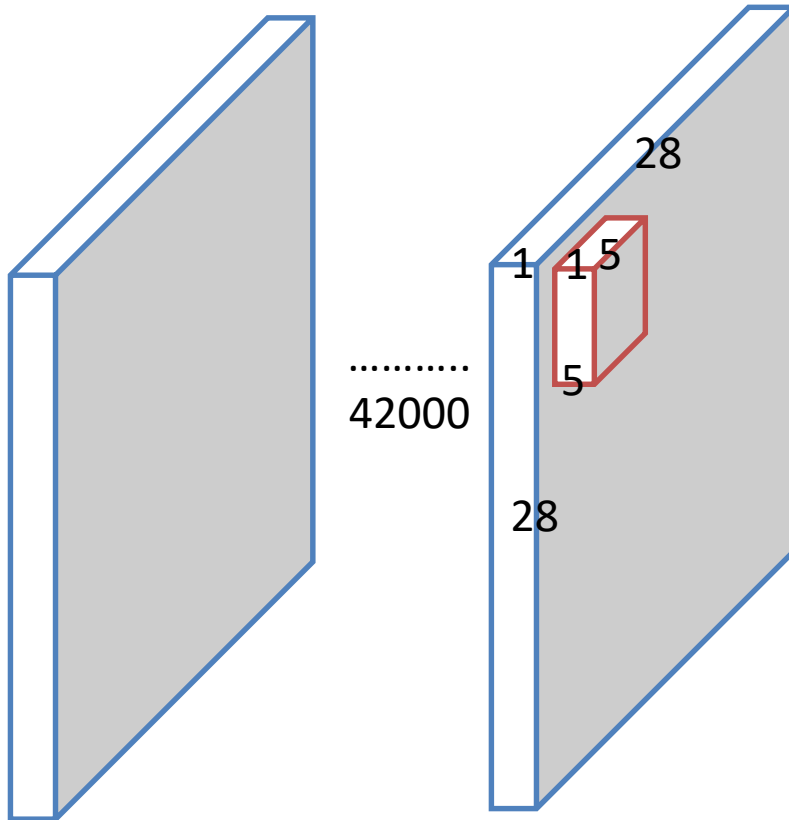
5x5x1



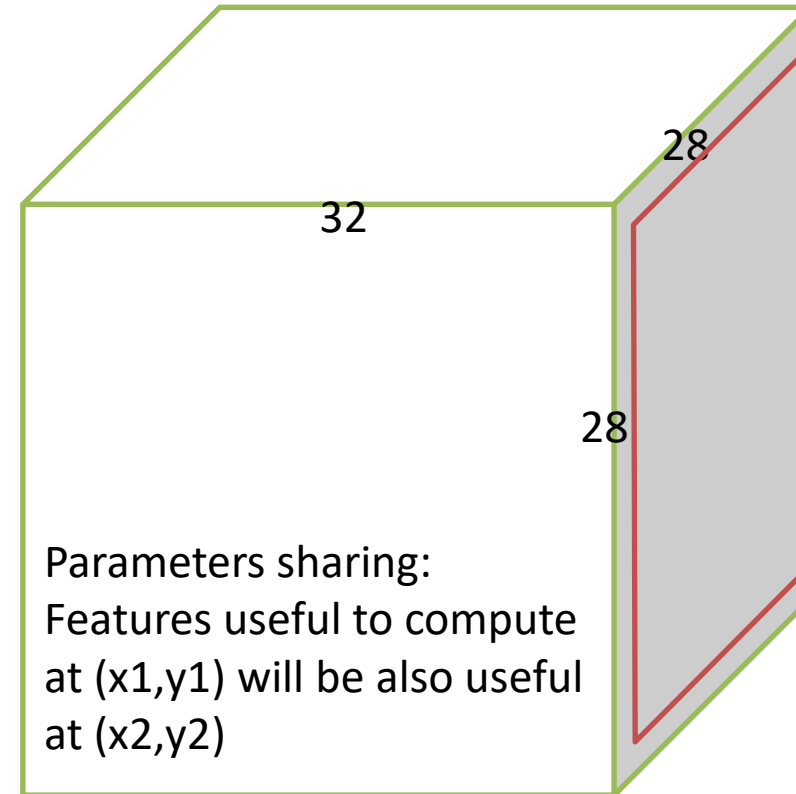
5x5x1

FC: 28x28x1 too large

Conv: connected to a small region of the layer



Padding = 'same'  
Equal input and  
output shape



28x28 each  
neurons  
calculate  
gradients



$\Sigma$   
Update a  
single  
5x5x1  
weights

28x28x32 neurons  
5x5x1 parameters

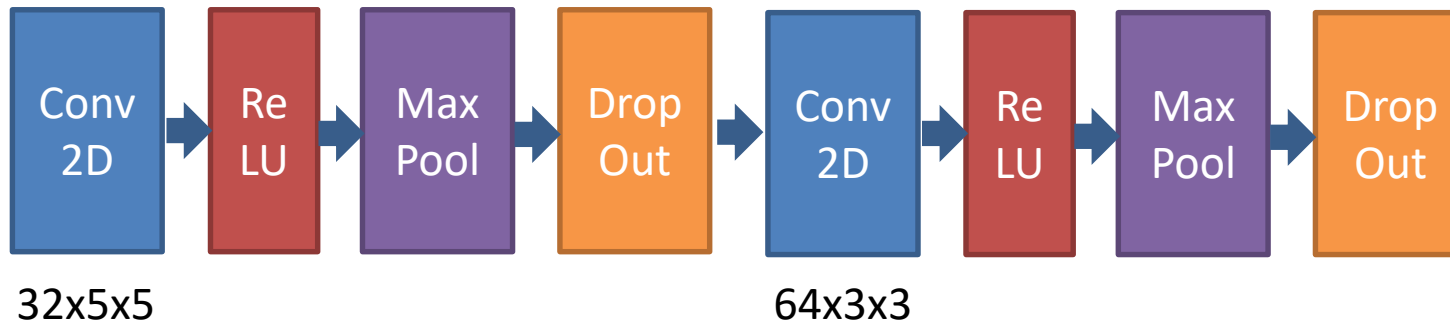


32x5x5x1 parameters

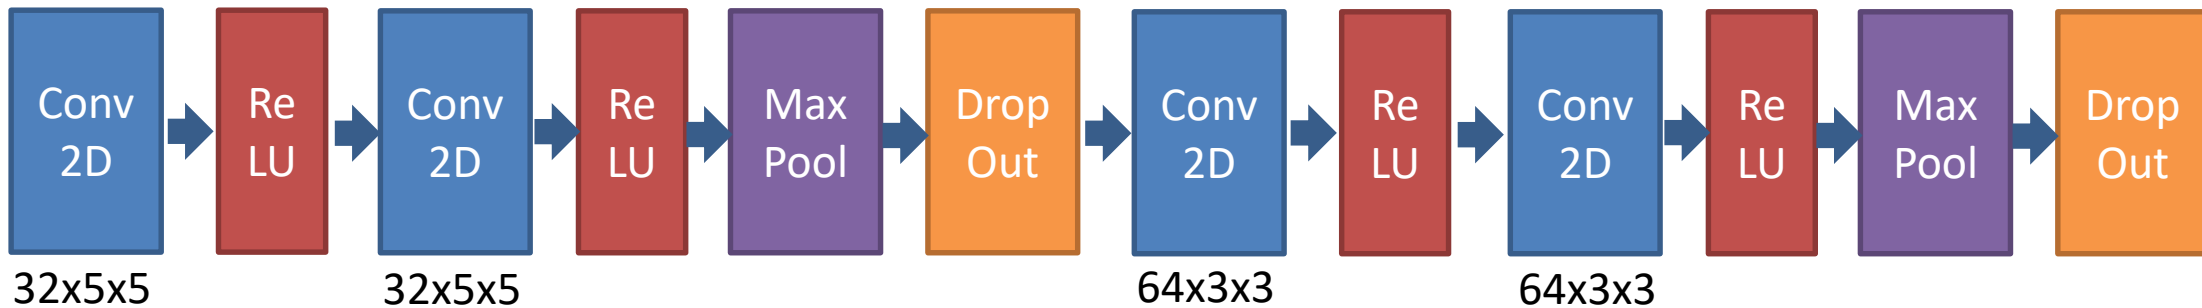
# Convolutional Layers

32x5x5 overall features, more large-scale features

64x3x3 local features, more detailed features, more filters

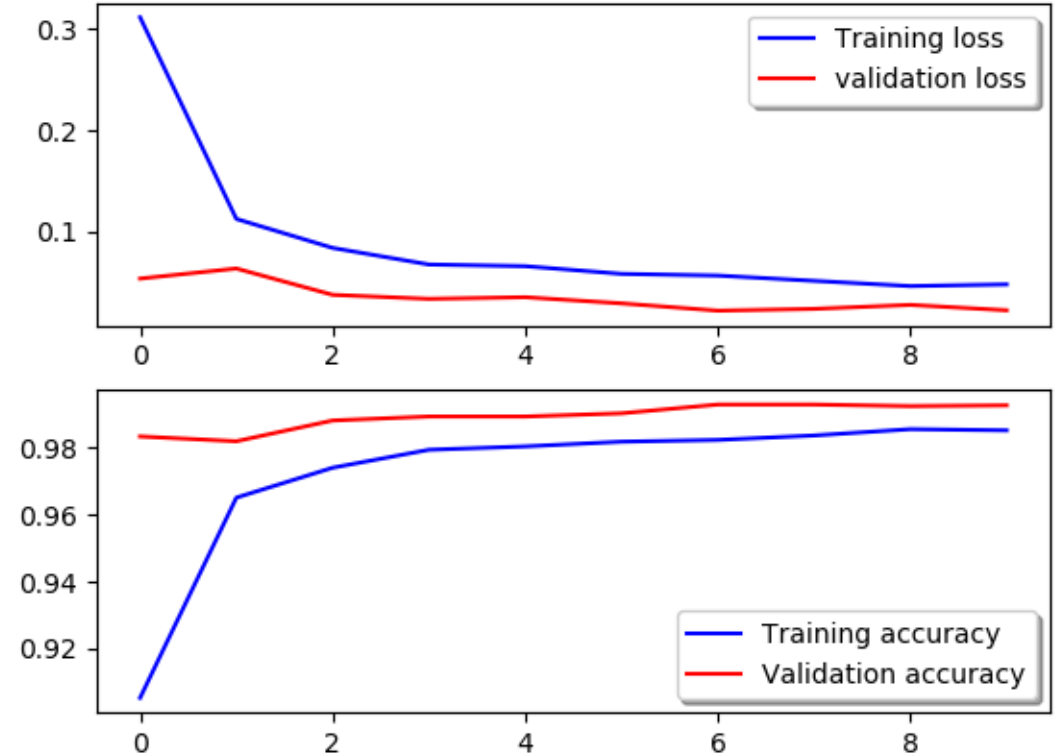
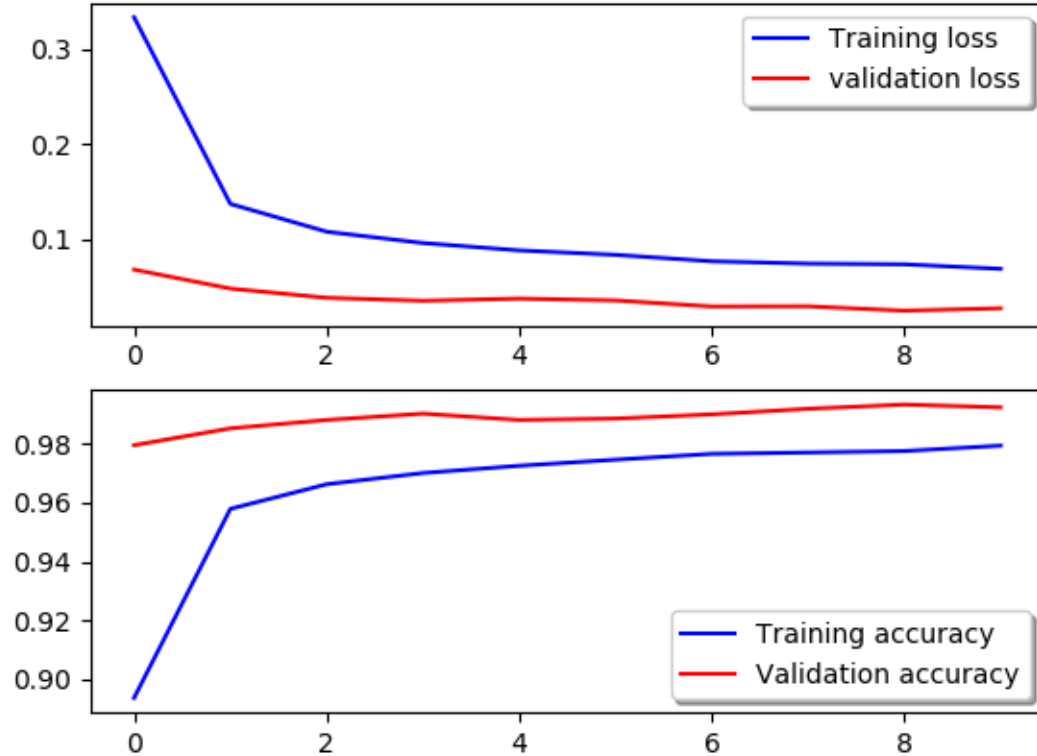


**Final loss: 0.0277**  
**Final accuracy: 0.9923**



**Final loss: 0.0226**  
**Final accuracy: 0.9926**

# Convolutional Layers



Training curves are closer to the validation curves --- more Convolutional layers tend to overfit more



# Data Argumentation

Reason for why validation loss curve is lower  
Training harder to identify

Not to overfit, artificially expand the dataset

Someone will write bigger/smaller numbers, scale

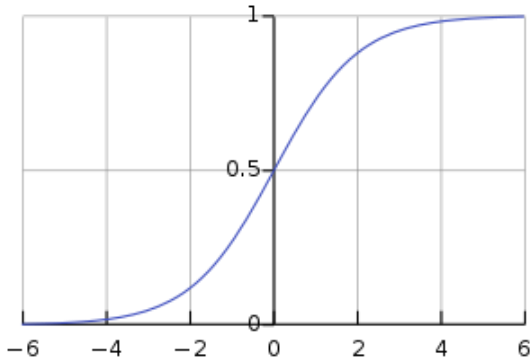
Rotate

Shift horizontally/vertically

```
datagen = ImageDataGenerator(  
    rotation_range=10,  
    zoom_range = 0.1,  
    width_shift_range=0.1,  
    height_shift_range=0.1)
```

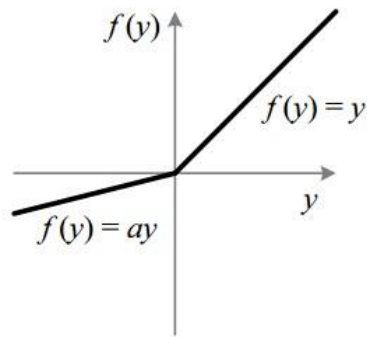
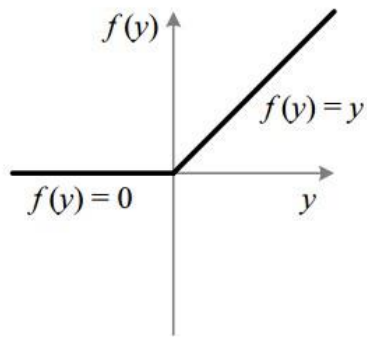
# Activation Layers

Add non-linearity to the models



Sigmoid & Tanh

- Gradient vanishing (derivative =  $g(x)(1-g(x))$ )
- Not zero-centred (always greater than 0, zig-zagging dynamics)
- Time-consuming to calculate



ReLU  $\max(0, x)$

- Not zero-centred
- Simply thresholding
- Can 'die': Every input  $x$  put ReLU zero  $\rightarrow$  die  
At least one  $x$  activate ReLU (fixed by small  $I_r$ )

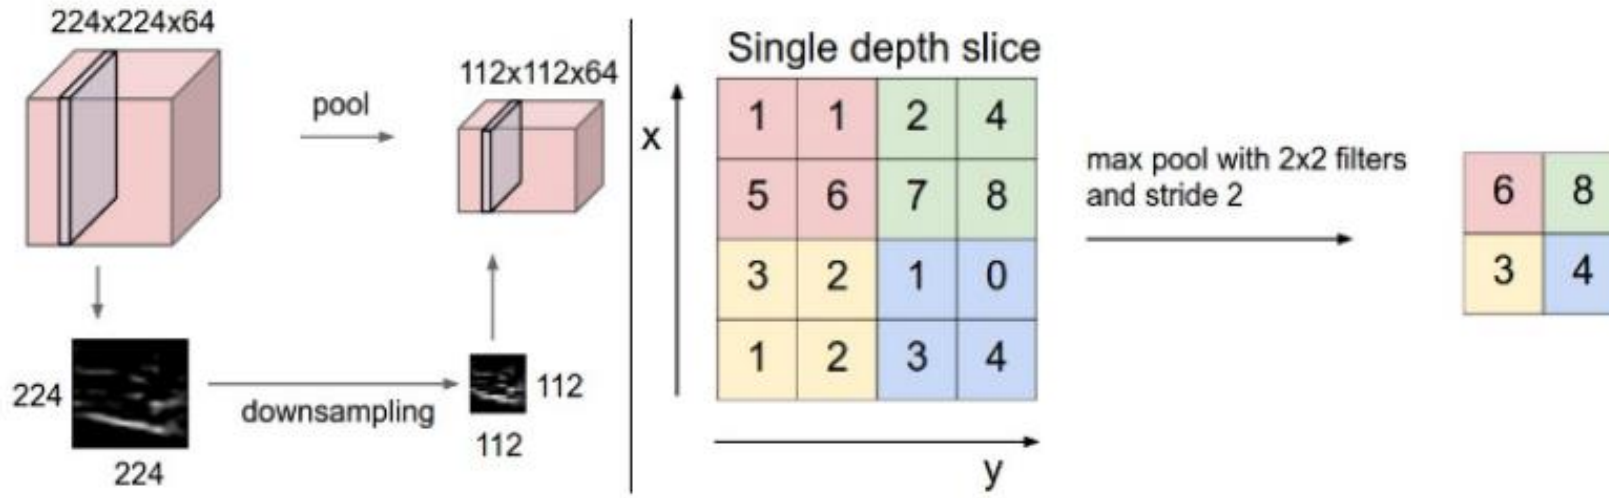
Leaky ReLU  $\max(0.01x, x)$

- Should be better since solved not zero-centred problem
- **Accuracy decrease from 0.9926 to 0.9864**
- **Error raise from 0.0226 to 0.0431**



# MaxPooling Layers

Down-sampling, pick maximum value in a 2x2 square  
Reduce computation cost, reduce overfitting (reduce dimension)  
 $28 \times 28 \rightarrow 14 \times 14$



# Dropout Layers

Randomly ignore some nodes in layer (Making new nns)  
Forces network to learn in a distributed way

$\approx$  training different neural network and then take average

# Batch Normalization Layers

- Internal covariance shift
- Change in the distribution of network activations due to the change in network parameters during training
- $F_2(F_1(x, \theta_1), \theta_2), \theta_2$  does not have to readjust to compensate for the change in the distribution of  $x$

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_1 \dots x_m\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

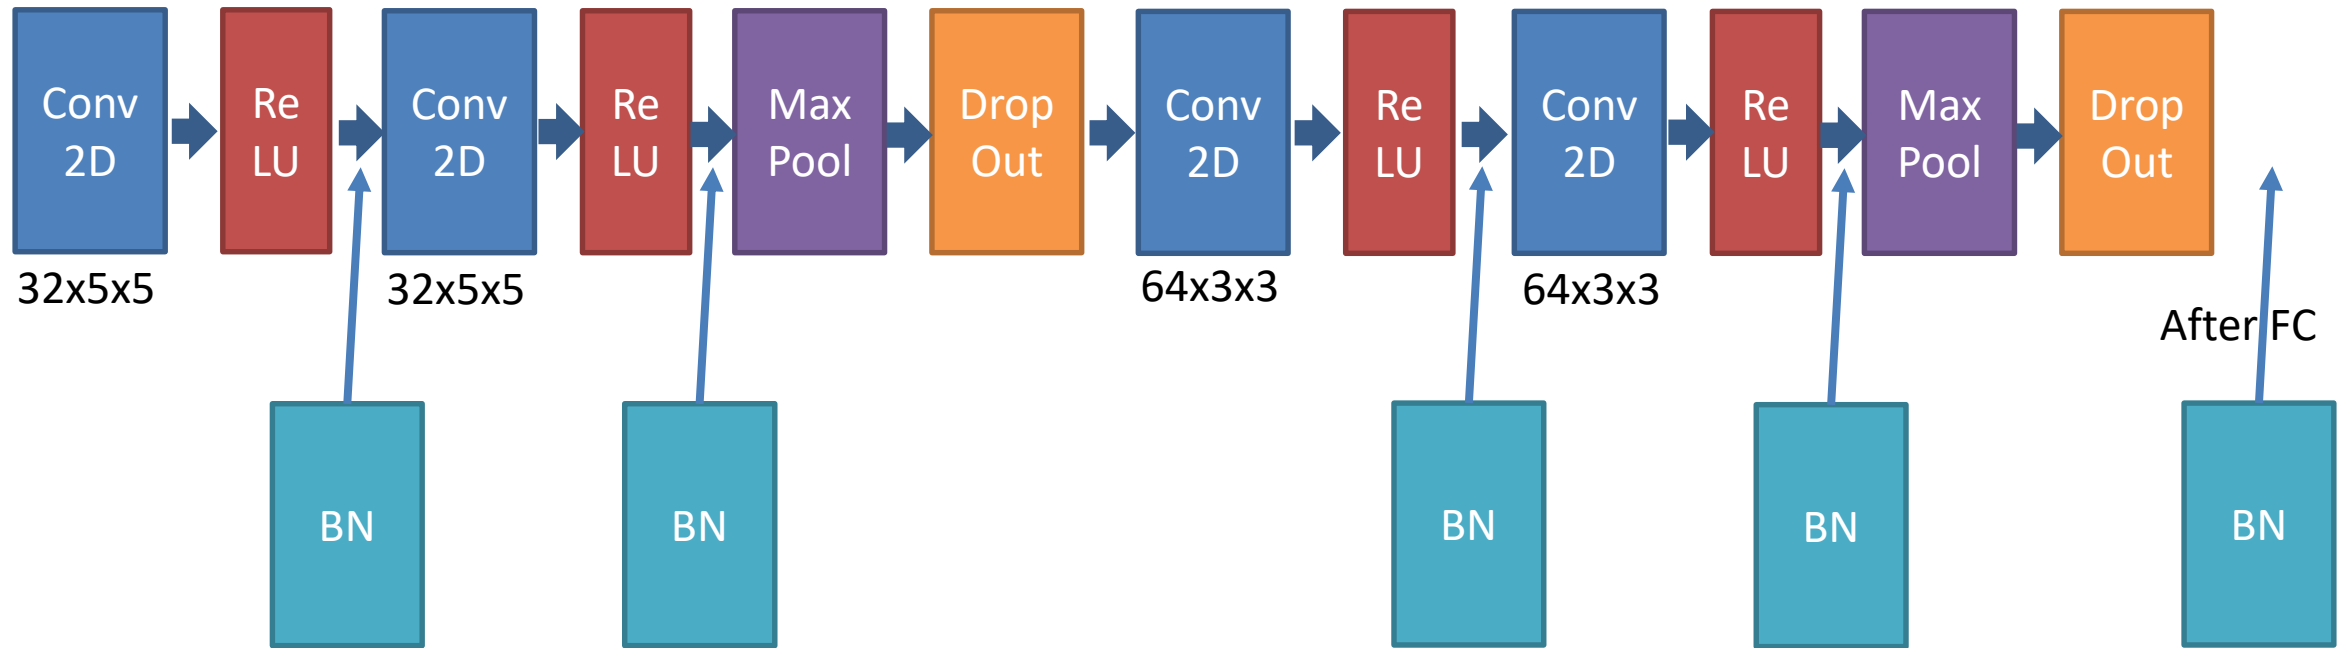
$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean} \quad \bullet \quad \text{Normalize each scalar feature independently}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance} \quad \bullet \quad \text{Mini-batch estimate mean \& variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift} \quad \bullet \quad \text{New parameters are learned, in order to restore the representation power of the network}$$

# Batch Normalization Layers



**4Conv2D + BN + ReLU**

**Accuracy: 0.9926    Error: 0.0226**

**4Con2D + ReLU**

**Accuracy: 0.9905    Error: 0.0323**

# Flatten + FC + Softmax Layers

- After Flatten: 1D vector
- After Last FC layer: 1D vector: 10x1x1 (10 categories)
- $\sigma(x) = \text{Softmax}(x) = \text{normalized}(\exp(x)) = \frac{\exp(x_i)}{\sum_j \exp(x_j)}$
- Endow score with meaning(probability distribution), Sum up to be 1
- Softmax + categorical\_crossentropy (loss function, categorical classification >2)
- When calculating gradient descent, derivative of softmax(x) =  $\sigma(x)(1-\sigma(x))$
- Get eliminated by the derivative of cross-entropy loss

# Optimizer

Functions to iteratively improve parameters

- SGD (slow), mini-batch Gradient Descent (partial)
  - Every time different batch of inputs (not stable)
- Momentum
  - Consider previous gradients

$$v_t = \gamma \cdot v_{t-1} + \alpha \cdot \nabla_{\Theta} J(\Theta)$$

$$\Theta = \Theta - v_t$$

- Adagrad

$$n_t = n_{t-1} + g_t^2$$

$$\Delta\theta_t = -\frac{\eta}{\sqrt{n_t + \epsilon}} * g_t$$

- RMSprop

**4Conv2D + ReLU + RMSprop**

**Accuracy: 0.9905    Error: 0.0323**

$$E[g^2]_t = 0.9E[g^2]_{t-1} + 0.1g_t^2$$

$$\Theta_{t+1} = \Theta_t - \frac{\alpha}{\sqrt{E[g^2]_t + \epsilon}} \cdot g_t$$

- Adam

**4Conv2D + ReLU + Adam**

**Accuracy: 0.9933    Error: 0.0232**

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$$

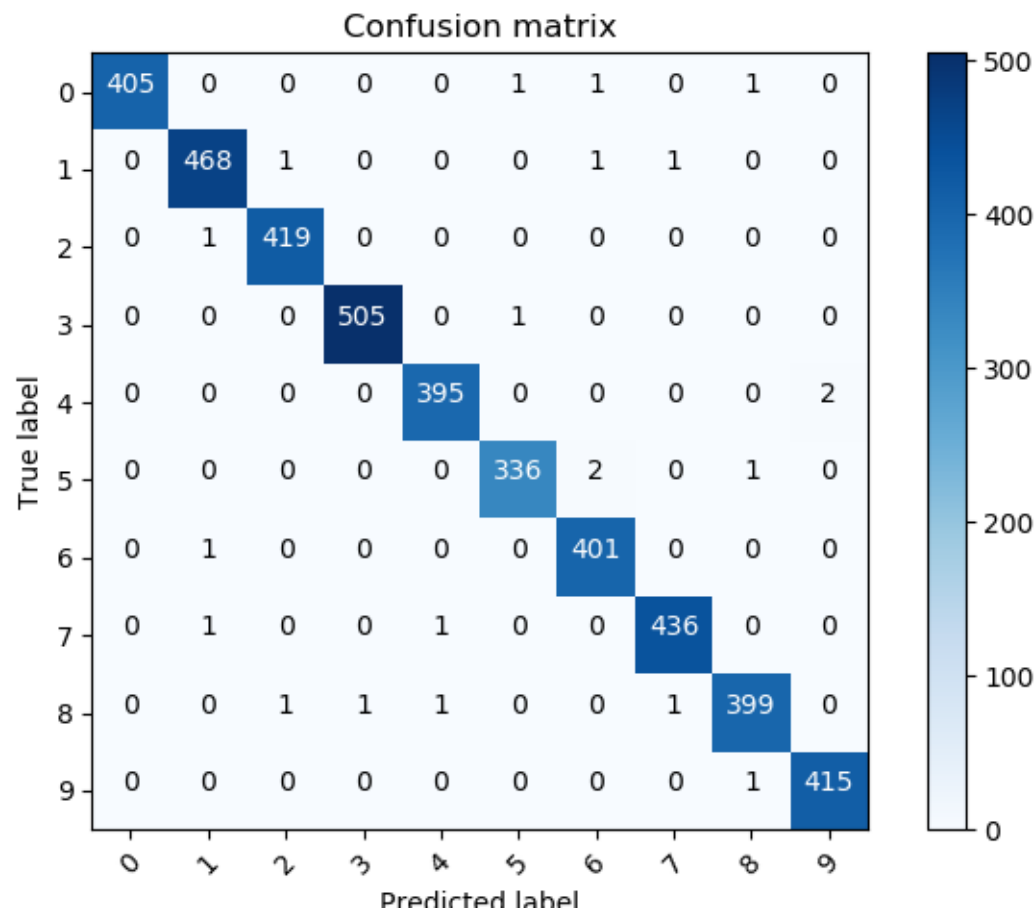
$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}$$

$$\hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

$$\Theta_{t+1} = \Theta_t - \frac{\alpha}{\sqrt{\hat{v}_t + \epsilon}} \hat{m}_t$$

# Result Evaluation : Confusion Matrix & Classification Report

[Conv2D + ReLU + BN] x2 + MaxPooling + DropOut + [Conv2D + ReLU + BN] x2 + MaxPooling + DropOut + Flatten + FC + DropOut + FC + Softmax  
Epochs → 30



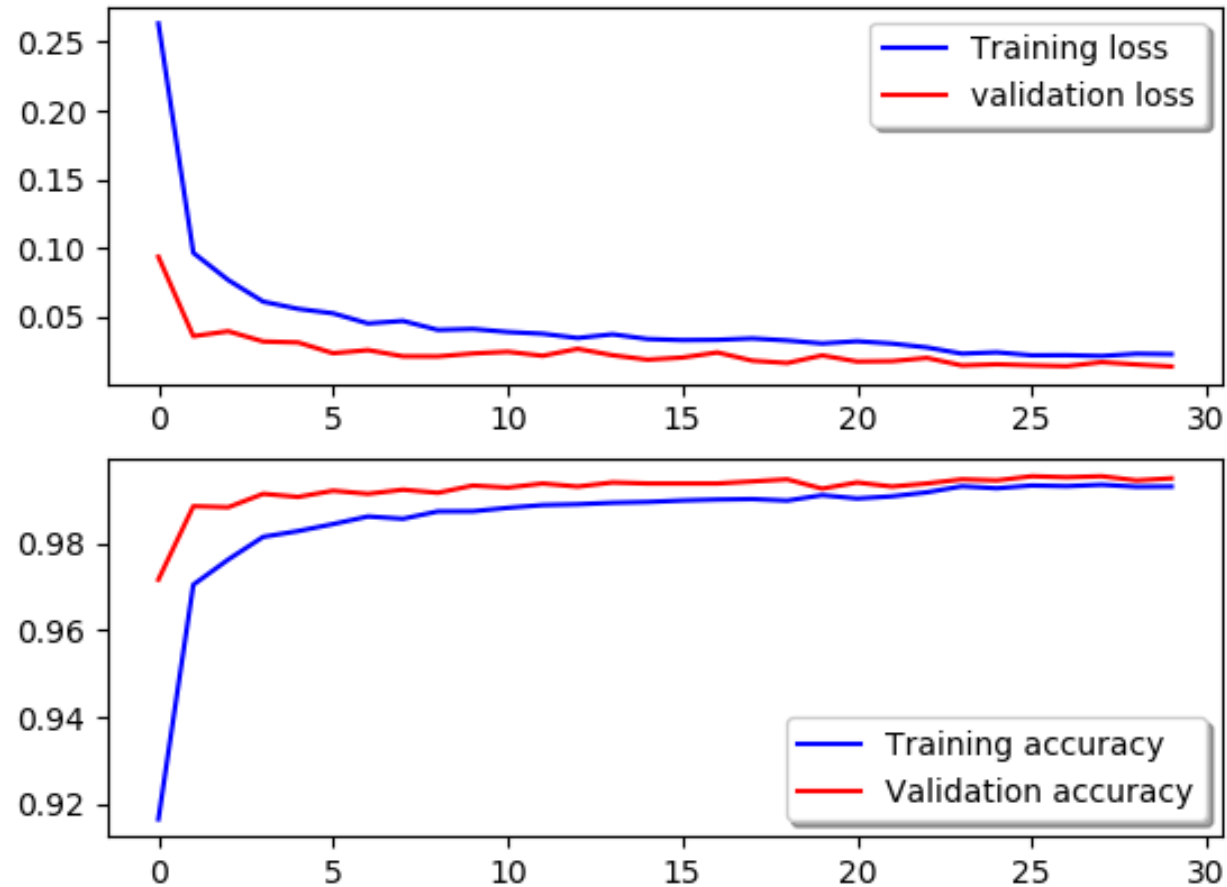
Final loss: 0.013653, final accuracy: 0.995000

	precision	recall	f1-score	support
0	1.00	0.99	1.00	408
1	0.99	0.99	0.99	471
2	1.00	1.00	1.00	420
3	1.00	1.00	1.00	506
4	0.99	0.99	0.99	397
5	0.99	0.99	0.99	339
6	0.99	1.00	0.99	402
7	1.00	1.00	1.00	438
8	0.99	0.99	0.99	403
9	1.00	1.00	1.00	416
avg / total	1.00	0.99	0.99	4200



# Result Evaluation : Training & Validation Curves

[Conv2D + ReLU + BN] x2 + MaxPooling + DropOut + [Conv2D + ReLU + BN] x2 + MaxPooling + DropOut + Flatten + FC + DropOut + FC + Softmax  
Epochs → 30



# Result Evaluation : Top 6 Errors

[Conv2D + ReLU + BN] x2 + MaxPooling + DropOut + [Conv2D + ReLU + BN] x2 + MaxPooling + DropOut + Flatten + FC + DropOut + FC + Softmax  
Epochs → 30

